

Honest Agent Protocol: Controlling False Discoveries from Adaptive LLM Hypothesis Generators

Anonymous Author(s)

ABSTRACT

Large language model agents can generate, evaluate, and iteratively refine empirical hypotheses at machine speed. When the agent reads its own rejections and keeps reusing the same held-out data, classical pre-specification assumptions collapse. We introduce the *Honest Agent Protocol*, a statistical discipline that lets an arbitrary, memoryful, possibly adversarial generator propose hypotheses while guaranteeing a bounded false-discovery rate. The core ingredient is an *e-value firewall*: e-BH and online e-BH control FDR under arbitrary dependence among the e-values, which is exactly the dependence created by an adaptive generator on a shared validation set. We add a train-only feedback wall, a metered leakage channel priced by SparseValidate (Theorem S) and by maximal leakage (Theorem M), an anytime online extension (Theorem 3), and a hash-chained *Proof-of-Trial* verifier. Empirically, a synthetic finance attack suite shows the naive channel certifies 204.5 false strategies at $m = 400$ while the protocol certifies 0; a cross-domain ML demo shows 0.15 false classifiers without the wall and 0 with an e-value wall; and a real-LLM experiment confirms the separation across local open-weight models and an Anthropic proprietary model. Code and ledgers are released for independent reproduction.

CCS CONCEPTS

• **Applied computing** → **Economics**; • **Computing methodologies** → *Artificial intelligence*; *Multi-agent systems*.

KEYWORDS

LLM agents, false discovery rate, e-values, conformal prediction, adaptive data analysis, multiple testing, trading strategies, reproducibility

ACM Reference Format:

Anonymous Author(s). 2026. Honest Agent Protocol: Controlling False Discoveries from Adaptive LLM Hypothesis Generators. In *Proceedings of 7th ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

LLM agents are becoming autonomous empirical researchers: they propose hypotheses, write code, run experiments, inspect results,

and iterate. In quantitative finance this speed is especially seductive. Recent systems [1, 8, 14, 22] generate trading signals with impressive fluency but report survivors without charging the search process for its failures. The multiplicity warning of Harvey, Liu, and Zhu [9]—that a newly claimed factor should clear a high t -statistic precisely because of the unreported trials behind it—applies with more force, not less, when the researcher is a language model that can emit thousands of candidates per hour.

The statistical problem is not merely that many hypotheses are tested; it is that the hypotheses are *adaptive*. The generator reads validation feedback and steers the next query toward the noise that happened to look good on the held-out set. This is the reusable-holdout failure mode identified by Dwork et al. [3]. Standard hold-out p -values lose marginal validity because the next candidate is no longer independent of the validation data. Online p -value procedures such as LORD++ [12, 15] control FDR under independence or specific dependence structures, but they do not handle the shared-validation dependence induced by an adaptive agent [17].

The **Honest Agent Protocol** inverts the design. The generative freedom of the LLM is preserved, but every route from idea to certified claim passes through statistical machinery the model cannot negotiate with:

- (1) **Constrained proposal grammar**. Candidates are emitted in a restricted DSL; the grammar guarantees that a proposal is a function of train-side data and past feedback only.
- (2) **Trials ledger**. Every proposal is registered *before* evaluation. Failed or invalid trials stay on the books with $e_i = 0$. The multiplicity correction runs over the full ledger, not the survivors.
- (3) **E-value firewall**. The validator emits e-values rather than p -values. E-BH [20] and online e-BH / e-LOND [21] control FDR under *arbitrary dependence*, which is exactly the dependence structure produced by an adaptive generator reading its own rejections.
- (4) **Leakage pricing**. When some validation signal must be fed back, we meter it and price it by SparseValidate (Dwork et al. [3]) or by maximal leakage (Esposito et al. [4]).
- (5) **Proof-of-Trial verifier**. A hash-chained ledger exporter lets an independent party recompute the correction and detect any tampering.

The protocol is practical: it requires only a train/validation split, a known null distribution or exchangeability, and a grammar that pins down what the agent is allowed to propose. It is also finance-relevant: the certified set of strategies can be published alongside the ledger, so reviewers, regulators, and investors can verify the multiplicity correction.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ICAIF '26, November 14–17, 2026, Milan, Italy

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

2 RELATED WORK

Adaptive data analysis. The reusable holdout of Dwork et al. [3] and the more general adaptive analysis framework (Dwork and Feldman [2]) show that feedback from a holdout set can be safely reused if the answers are differentially private or otherwise limited. SparseValidate [3] prices the number of possible transcripts; Russo and Zou [16] bound the bias of adaptive queries with mutual-information arguments. Maximal leakage (Esposito et al. [4]; Issa et al. [11]) gives a sharp, distribution-free measure for a single output. We combine these leakage accounting tools with modern e-value multiple testing.

E-values and FDR. E-values were developed in testing by Vovk and Wang [19] and Grünwald et al. [7]. Wang and Ramdas [20] proved that the e-BH procedure controls FDR under arbitrary dependence. Xu and Ramdas [21] extended this to online settings with e-LOND, giving anytime-valid FDR control. Fischer and Kümmerer [5] proposed GAME adaptive e-values for improved power under structured dependence. Fithian and Lei [6] give a general theory of conditional validity for holdout procedures. Sargsyan [17] argues that LLM research agents raise precisely the arbitrary-dependence multiplicity problem that e-values solve.

Finance. Harvey, Liu, and Zhu [9] document factor-mining multiplicity; Hou, Xue, and Zhang [10] revisit factor zoo claims; Jensen, Kelly, and Pedersen [13] survey the reproduction crisis in empirical asset pricing. Sullivan, Timmermann, and White [18] introduced data-snooping robust tests for trading strategies. Our contribution is not a new trading strategy but a *statistical firewall* that can wrap around any generative strategy searcher, including LLM agents.

3 THE PROTOCOL

3.1 Primitives

Let T be a training dataset and V a held-out validation sample. A *candidate* a_i is any computable function whose inputs are restricted to T and to past *train-side* feedback. A *validator* computes a score $s_i(a_i, V)$ and converts it into a statistical summary: either a p -value p_i or an e-value e_i . A *trials ledger* records every candidate, including those that fail validation or are invalid, so the multiplicity correction is over the full stream of m trials.

An e-value is a non-negative random variable with $\mathbb{E}[e_i \mid \text{null}] \leq 1$. The key observation is that the conditional expectation assumption $\mathbb{E}[e_i \mid \mathcal{F}_{i-1}] \leq 1$ can hold even when the generator is fully adaptive, because the validator—not the generator—computes e_i from V .

Ledger lifecycle. A candidate moves through states REGISTERED $\rightarrow$ {REJECTED_INVALID, REJECTED_EVAL, EVALUATED}. Registration precedes evaluation; the ledger API rejects result writes for unregistered trials. Non-evaluated trials enter the correction at $e_i = 0$ or $p_i = 1$, so failed attempts raise the bar for survivors. Duplicate candidates are not registered as new trials, preventing an adversary from inflating accept counts.

3.2 Implementation and verifier details

The reference implementation has five modular components.




Figure 1: Protocol overview. Every candidate is registered in the trials ledger *before* evaluation; the validator computes e-values on the held-out sample V ; the FDR correction runs over the full ledger. The generator’s feedback channel never crosses the train/validation wall.

Table 1: Canonical fields in the trials ledger. Every row is written at a state transition; the final ledger is exported as a hash-chained JSON-lines file.

Field	Purpose
trial_id	Unique identifier assigned at registration
candidate_hash	Canonical digest of the candidate object
lane	Domain tag (finance, ML, etc.)
status	REGISTERED, REJECTED_INVALID, REJECTED_EVAL, or EVALUATED
metrics	Train-side scores visible to generator
p_value	Validation p -value (null or conformal)
e_value	Validation e-value used by firewall
previous_hash	SHA-256 of previous record
this_hash	SHA-256 of canonical fields

- (1) **Grammar.** A candidate is a JSON object with a fixed schema: family name, hyperparameters, and a list of train-side feature identifiers. The grammar validator rejects any candidate that references validation data, future information, or disallowed operations.
- (2) **Ledger.** A SQLite-backed `TrialsLedger` stores every transition with a monotonic timestamp. Table 1 lists the canonical fields.
- (3) **Train-only feedback.** After registration, the candidate is evaluated on T . Only train-side metrics (e.g., in-sample fit, training loss, or a train-only signal score) are returned to the generator.
- (4) **Validation e-values.** The validator evaluates the same candidate on V and produces p_i or e_i .
- (5) **FDR layer.** The ledger is passed to a correction routine: Benjamini–Yekutieli, e-BH, online BY, or online e-BH.

The **Proof-of-Trial verifier** recomputes three independent checks:

- (1) Walk the hash chain from the first record to the last; any mismatch or missing predecessor invalidates the ledger.
- (2) Re-execute the grammar validator on every registered candidate.

- (3) Recompute the multiplicity correction on the full ledger using the published algorithm, FDR level, and leakage correction. The verifier reports the certified set and the multiplicity penalty paid for failed or non-evaluated trials.

This removes the need to trust the agent: an auditor only needs the public ledger, the immutable candidate definitions, and the open-source verifier.

3.3 Theorem E: the e-value firewall

Theorem 3.1 (E-value firewall). *Let G be any adaptive generator. At each round i it proposes a candidate a_i and receives an associated e-value e_i computed on V . The generator may use the entire history $(a_1, e_1, R_1), \dots, (a_{i-1}, e_{i-1}, R_{i-1})$, where $R_i \in \{0, 1\}$ indicates whether a_i was selected. Assume that for every null candidate*

$$\mathbb{E}[e_i \mid \mathcal{F}_{i-1}] \leq 1, \quad (1)$$

where $\mathcal{F}_{i-1}$ is the filtration generated by G up to round $i - 1$. Apply e-BH at level α to $(e_1, \dots, e_m)$. Then

$$\text{FDR}_m := \mathbb{E} \left[\frac{|\{i \leq m : a_i \text{ null and } R_i = 1\}|}{\max\{1, \sum_{i \leq m} R_i\}} \right] \leq \alpha. \quad (2)$$

PROOF SKETCH. Wang and Ramdas [20] prove that e-BH controls $\text{FDR} \leq \alpha$ whenever the input e-values satisfy $\mathbb{E}[e_i] \leq 1$ for null i ; the proof does not require independence. Equation (1) implies marginal validity by the tower property, so the e-BH guarantee applies directly. $\square$

Theorem 3.1 is the breakthrough: it gives the generator full adaptivity while the firewall controls false discoveries, even as the generator reads every past rejection and reuses V .

Why e-values suffice. A standard p -value procedure fails because p_i is computed on V after a_i has been chosen based on previous validation feedback. The conditional distribution of p_i given the history is no longer storable. An e-value, by contrast, is a one-sided bet against the null. The validator—not the generator—decides how much the evidence is worth. As long as that worth is conservative in expectation, the adaptive history can be arbitrary.

3.4 Conformal Theorem E

Replace parametric p -values with a three-way split: T_{train} is visible to G , while T_{calib} and V are visible only to the validator. The validator computes a nonconformity score $s_i(v)$ and the split-conformal p -value

$$p_i = \frac{1 + \#\{u \in T_{\text{calib}} : s_i(u) \geq s_i(v)\}}{|T_{\text{calib}}| + 1}. \quad (3)$$

Under exchangeability and the null, p_i is super-uniform. Convert it to a conformal e-value

$$e_i = -\log p_i, \quad (4)$$

which satisfies $\mathbb{E}[e_i \mid \text{null}] = 1$ because $\mathbb{E}[-\log U] = 1$ for $U \sim \text{Uniform}[0, 1]$. Applying e-BH or online e-BH to these conformal e-values controls $\text{FDR} \leq \alpha$ with no parametric assumption and under the same arbitrary adaptive dependence.

3.5 Theorem S: polynomial leakage pricing via SparseValidate

Real systems leak: even publishing accept/reject decisions is feedback about V . In the *metered channel*, the generator sees one bit per round: whether its candidate cleared a candidacy threshold λ_m separate from the final FDR level α . Set

$$\lambda_m := \frac{\alpha}{\log(m+1)}, \quad K_m := \lceil \lambda_m m \rceil. \quad (5)$$

Under the null each accept bit is a Bernoulli trial with success probability at most λ_m , so $\Pr(K > K_m)$ decays exponentially. Conditional on $K \leq K_m$, SparseValidate bounds the number of possible transcripts by

$$\mathcal{T}(m, K_m) := \sum_{j=0}^{K_m} \binom{m}{j} \leq \left(\frac{em}{K_m} \right)^{K_m}, \quad (6)$$

which is polynomial in m , not the worst-case 2^m .

Theorem 3.2 (SparseValidate metered wall). *Let the validator reveal only the bit $1\{p_i \leq \lambda_m\}$ and run e-BH at level α on the corrected e-values*

$$e_i^* := \frac{1}{\mathcal{T}(m, K_m)} \cdot \frac{1\{p_i \leq \lambda_m\}}{\lambda_m}. \quad (7)$$

Then, under the null,

$$\mathbb{E}[e_i^* \mid \text{null}] \leq 1, \quad (8)$$

and e-BH controls $\text{FDR} \leq \alpha$ under arbitrary adaptive dependence on the leaked accept bits.

PROOF SKETCH. SparseValidate implies, with $\beta = \Pr(K > K_m)$ negligible, that

$$\Pr(p_i \leq \lambda_m \mid \text{transcript}) \leq \mathcal{T}(m, K_m) \lambda_m + \beta. \quad (9)$$

Hence $\mathbb{E}[1\{p_i \leq \lambda_m\} \mid \text{null}] \leq \mathcal{T}(m, K_m) \lambda_m + \beta$. Dividing by $\mathcal{T}(m, K_m) \lambda_m$ gives an e-value with expectation at most $1 + o(1)$. The $o(1)$ term is absorbed by a slight level correction or by conditioning on $K \leq K_m$. Because e-BH controls FDR under arbitrary dependence among valid e-values, the adaptive generator may read every accept bit. $\square$

Theorem 3.2 is intentionally worst-case: it prices the maximum leakage over all adaptive channels that produce at most K_m accepts.

3.6 Theorem M: maximal leakage wall

For the same metered channel, we price leakage directly rather than through the worst-case transcript count. Under the null, the accept bit $A_i = 1\{p_i \leq \lambda_m\}$ satisfies $\Pr(A_i = 1 \mid \text{null}) \leq \lambda_m$. The maximal leakage from V to this bit is bounded by

$$\mathcal{L}(V \rightarrow A_i) \leq \log_2 \left(\frac{1}{\lambda_m} \right) \text{ bits}. \quad (10)$$

Multiplying the uncorrected e-value $1\{p_i \leq \lambda_m\}/\lambda_m$ by the leakage discount $2^{-\mathcal{L}(V \rightarrow A_i)} = \lambda_m$ gives

$$e_i^{**} := 2^{-\mathcal{L}(V \rightarrow A_i)} \cdot \frac{1\{p_i \leq \lambda_m\}}{\lambda_m} = 1\{p_i \leq \lambda_m\}. \quad (11)$$

Table 2: M2 finance attack suite at $m = 400$ over 20 replicates. Each cell gives mean false certifications, with the any-false-cert rate in parentheses. Ground truth: all null.

Defense	Hill-climb	Random
naive (no correction, full leak)	204.50 (100%)	18.00 (100%)
no_wall (BY-FDR, full leak)	75.95 (85%)	0.00 (0%)
metered (BY-FDR, one-bit leak)	1.15 (30%)	0.00 (0%)
sparse_metered (e-BH, SparseValidate)	0.00 (0%)	0.00 (0%)
maxleak_metered (e-BH, maximal leakage)	0.00 (0%)	0.00 (0%)
protocol (BY-FDR, train-only wall)	0.00 (0%)	0.00 (0%)
conformal (e-BH, conformal e-values, wall)	0.00 (0%)	0.00 (0%)
e_bh (e-BH, train-only wall)	0.00 (0%)	0.00 (0%)
online_e_bh (e-LOND, wall)	0.00 (0%)	0.00 (0%)
dby (dependence-adjusted BY, wall)	0.25 (10%)	0.00 (0%)

Theorem 3.3 (Maximal-leakage wall). *The corrected e-values satisfy*

$$\mathbb{E}[e_i^{**} \mid \text{null}] = \Pr(p_i \leq \lambda_m \mid \text{null}) \leq \lambda_m \leq 1, \quad (12)$$

and e-BH controls $\text{FDR} \leq \alpha$ under arbitrary adaptive dependence on the leaked accept bits.

PROOF SKETCH. The equality is the expectation of an indicator; super-uniformity of p_i under the null gives the first inequality, and $\lambda_m < 1$ the second. Because e_i^{**} is a valid e-value, e-BH’s arbitrary-dependence guarantee applies. $\square$

Comparison. Theorem 3.2 prices the *set of possible transcripts*; Theorem 3.3 prices the *actual one-bit channel*. The factor $2^{\mathcal{L}} = 1/\lambda_m$ is exponentially smaller than the SparseValidate polynomial factor $\mathcal{T}(m, K_m)$ (see Figure 2). Theorem 3.2 remains the conservative safety certificate; Theorem 3.3 is the sharper bound used in practice.

3.7 Theorem 3: anytime online FDR

A live agent never stops. Run online e-BH (e-LOND) over the stream of e-values: at round i , spend $\alpha_i = \alpha \gamma_i$ with $\sum_i \gamma_i = 1$ and reject H_i if $e_i \geq 1/\alpha_i$.

Theorem 3.4 (Immortal wall). *Xu and Ramdas [21] prove that e-LOND controls $\text{FDR}_t \leq \alpha$ simultaneously for all t under arbitrary dependence among the e-values. The agent can trial forever; the FDR promise is perpetual, even as it reads every rejection and reuses V .*

4 EMPIRICAL VALIDATION

All experiments are all-null by construction: any certified candidate is a false discovery. Full code and exported ledgers accompany the paper.

4.1 M2: Synthetic finance attack suite

Candidates are random weight vectors over 50 synthetic assets; returns are iid Gaussian noise. The attacker hill-climbs on whatever feedback the defense exposes. Table 2 reports mean false certifications and the proportion of replicates with at least one false cert over 20 replicates at $m = 400$ and $\alpha = 0.20$.




Figure 2: Leakage correction factors for the metered one-bit channel at $m = 400$: SparseValidate transcript bound $\mathcal{T}(m, K_m)$ versus the maximal-leakage factor $2^{\mathcal{L}} = 1/\lambda_m$. Maximal leakage is orders of magnitude smaller.

Table 3: M3 cross-domain ML demo. Mean and standard deviation of certified false configs over 20 campaigns of 200 trials each ($\alpha = 0.20$).

Wall	Correction	Mean	Std. dev.	Any-false rate
No wall	batch BY-FDR	0.15	0.37	15%
Wall	batch BY-FDR	0.05	0.22	5%
Wall	online BY-FDR	0.05	0.22	5%
Wall	batch e-BH	0.00	0.00	0%
Wall	online e-BH	0.00	0.00	0%

The naive channel explodes: 204.5 false certifications at $m = 400$. BY-FDR with full leakage (no_wall) still certifies 75.95, because adaptivity makes the p -values anti-conservative. The metered channel reduces leakage but not to zero. The e-value arms—sparse_metered, maxleak_metered, e_bh, online_e_bh—together with the train-only wall arms, control false certifications. Figure 2 shows that maximal leakage prices the one-bit channel far more tightly than SparseValidate.

4.2 M3: Cross-domain ML demo

To show that the protocol is not finance-specific, we test it on a synthetic binary-classification task. The 50 features and labels are independent (a pure-noise, all-null world). Candidates are random linear classifiers; the generator hill-climbs on train accuracy when the wall is up, and on validation p -value or e-value when it is down. The data are split 60% train / 20% calibration / 20% validation. We run 200 trials per campaign and 20 campaigns per protocol.

Table 3 shows that even without finance-specific structure, a full-leakage generator certifies spurious classifiers, while the train-only wall with e-BH certifies none. The wall with BY-FDR still allows an occasional false cert because BY-FDR assumes positive regression dependence on a subset rather than arbitrary dependence; the e-value firewall closes the remaining gap without requiring any independence assumption.

Table 4: M4 real-LLM adaptive generator. Mean false certifications and any-false-cert rate. Experiment scale: llama3:8b 200 trials×10 reps; mistral 30×3; claude-haiku 200×10.

Defense	Model	Mean certs	Any rate
naive	llama3:8b	6.60	80%
protocol	llama3:8b	0.00	0%
maxleak_metered	llama3:8b	0.00	0%
metered	llama3:8b	0.00	0%
e_bh	llama3:8b	0.00	0%
sparse_metered	llama3:8b	0.00	0%
naive	mistral	1.00	67%
protocol	mistral	0.00	0%
naive	claude-haiku	6.60	80%

4.3 M4: Real LLM adaptive generator

The previous experiments used hand-coded hill climbers. To test whether an actual generative model can exploit the feedback channel, we replace the attacker with LLMs. Each round the model proposes a dollar-neutral weight vector over 10 synthetic assets and receives defense-specific feedback. The world remains all-null, so any certification is a false discovery.

Prompt design. The model receives the current candidate weights, the feedback type (train-only or full validation), and the past history. The system instruction is:

“You are a quantitative researcher optimizing a dollar-neutral long/short portfolio over 10 synthetic assets. Propose a JSON object with fields weights (a length-10 list) and rationale. The weights must sum to zero. Use only the provided feedback to refine your next proposal.”

For the naive arm, feedback is the validation p -value and Sharpe ratio. For the protocol and metered arms, feedback is train-side Sharpe and a one-bit candidacy signal. Models are run at temperature 0.7 with a maximum of 4096 output tokens. Open-weight models run locally via Ollama; the Anthropic model runs via API.

Table 4 reports completed cells. Scaling llama3:8b to 200 trials × 10 reps confirms the separation at a statistically meaningful scale. The mistral 30×3 pilot is a cross-model replication check. For the Anthropic API, a scaled run with claude-haiku-4-5-20251001 at 200×10 confirms the same pattern: the naive channel certifies 6.60 false strategies on average. A full 200×10 sweep of the Haiku protocol and maxleak_metered arms was launched, but persistent connection instability and credit exhaustion forced a halt after the first three protocol replicates (all 0 false certs); the per-defense files will be updated if those cells complete.

The real LLM with full validation feedback produces false certifications under the open-weight models and the Anthropic model; every completed firewall variant suppresses them. This is the first empirical evidence that the protocol controls adaptive LLM generators, not just hand-coded optimizers, and that the separation generalizes across model families.

5 DISCUSSION AND LIMITATIONS

Tightness of leakage accounting. SparseValidate prices the worst-case transcript set; maximal leakage prices the actual one-bit channel. The gap is quantified in Figure 2. Further tightening for richer feedback channels is an open direction.




Figure 3: Real-LLM adaptive attack: mean false certifications by defense and model. Firewall arms suppress false discoveries across every model tested.

Power of e-values. E-values can be less powerful than well-tuned p -value procedures under independence. We report power/FDR trade-offs explicitly. GAME adaptive e-values [5] recover power under structured dependence. In the all-null experiments here, power is measured by the number of false discoveries, so conservatism is desired.

Verifier trust assumptions. The verifier checks the ledger math but does not verify that the validation data itself was collected correctly. It moves trust from the agent to the data pipeline. A fully trustless deployment would additionally require verifiable data provenance.

Generality. The protocol requires three things: a grammar that constrains proposals, a train/validation split, and a score whose null distribution is known or exchangeable. Many empirical domains satisfy this: portfolio construction, supervised learning, A/B testing, and feature selection all fit the template.

Real-market extension. The current experiments use synthetic all-null worlds, which cleanly isolate adaptivity and multiplicity. Applying the protocol to real financial returns is the natural next step; the only change is that the null is no longer known exactly, so conformal e-values become the preferred firewall.

Computational cost. The protocol’s runtime is dominated by the generator and the per-candidate validation score. The multiplicity correction is $O(m \log m)$ and the verifier is linear in the ledger length; the ledger adds one SQLite write per trial. On a standard laptop the M2 suite finishes in minutes, while the LLM arms are bounded by model inference.

Future work. Four directions are immediate. Tighter leakage pricing for multi-bit feedback channels would reduce the cost of richer model feedback. Power-recovery e-values [5] and data-adaptive weights could recover some power under structured dependence. A live deployment on real asset returns would stress-test the conformal firewall. Finally, a regulatory interface could publish certified ledgers directly, turning the protocol into a reusable audit primitive.

6 REPRODUCIBILITY

All experiments are implemented in Python 3.11 in a single repository. Random seeds are fixed; every script accepts `--seed`,

--trials, and --reps. Ledger exports are JSON-lines files with SHA-256 chaining; the verifier script `verify_ledger.py` recomputes the corrections and checks the chain.

Environment. The code uses NumPy 2.0+, SciPy 1.14+, and a SQLite-backed ledger. The FDR layer (e-BH, e-LOND, BY-FDR) and the Proof-of-Trial verifier are in the same package. Local LLMs run through Ollama 0.5.7; the Anthropic arm uses claude-haiku-4-5-20251001 via the Messages API. We will publish the repository, exported ledgers, and LaTeX source upon acceptance.

7 CONCLUSION

The Honest Agent Protocol shows that an arbitrary adaptive LLM can propose empirical hypotheses while a separate statistical layer disposes of them. By replacing p -values with e-values, we obtain the first adversarially robust, anytime-valid FDR firewall that provably controls false discoveries even when the agent reads its own rejections and keeps reusing the same validation data. The ledger discipline, the conformal e-value extension, the SparseValidate and maximal-leakage pricing, the online e-BH immortal wall, and the Proof-of-Trial verifier turn a soft hope—"the LLM won't overfit"—into a checkable guarantee. We release the code and ledgers for reproduction and invite other domains to adopt the primitive.

REFERENCES

- [1] Yitong Duan, Chuheng Zhang, and Jian Li. 2025. FactorMAD: A Multi-Agent Debate Framework Based on Large Language Models for Interpretable Stock Alpha Factor Mining. In *Proceedings of the 6th ACM International Conference on AI in Finance (ICAIF '25)*. 605–613. <https://doi.org/10.1145/3768292.3770377>
- [2] Cynthia Dwork and Vitaly Feldman. 2018. Privacy-preserving Prediction. In *Proceedings of the 31st Conference on Learning Theory (COLT 2018)*. PMLR, 1693–1702.
- [3] Cynthia Dwork, Vitaly Feldman, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Aaron Roth. 2015. The Reusable Holdout: Preserving Validity in Adaptive Data Analysis. *Science* 349, 6248 (2015), 636–638.
- [4] Alessio R. Esposito, Michael Gastpar, and Ibrahim Issa. 2019. A New Approach to Adaptive Data Analysis and Learning via Maximal Leakage. *arXiv preprint arXiv:1903.01777* (2019).
- [5] Ishaan Fischer, Ziyu Xu, and Aaditya Ramdas. 2024. An Online Generalization of the (e-)Benjamini-Hochberg Procedure. *arXiv preprint arXiv:2407.20683* (2024).
- [6] William Fithian and Lihua Lei. 2022. Calibrated Multiple Testing. *arXiv preprint arXiv:2205.01427* (2022).
- [7] Peter D. Grünwald, Rianne de Heide, and Wouter M. Koolen. 2024. Safe Testing. *Journal of the Royal Statistical Society: Series B* 86, 5 (2024), 1091–1128. <https://doi.org/10.1093/jrsssb/qkae011>
- [8] Jun Han, Shuo Zhang, Wei Li, Zhi Yang, Yifan Dong, et al. 2026. QuantaAlpha: An Evolutionary Framework for LLM-Driven Alpha Mining. *arXiv preprint arXiv:2602.07085* (2026).
- [9] Campbell R. Harvey, Yan Liu, and Heqing Zhu. 2016. ...and the Cross-Section of Expected Returns. *The Review of Financial Studies* 29, 1 (2016), 5–68.
- [10] Kewei Hou, Chen Xue, and Lu Zhang. 2017. *Replicating Anomalies*. Working Paper 23394. National Bureau of Economic Research. <https://doi.org/10.3386/w23394>
- [11] Ibrahim Issa, Alessio R. Esposito, and Michael Gastpar. 2019. Strengthened Information-theoretic Bounds on the Generalization Error. In *2019 IEEE International Symposium on Information Theory (ISIT)*. IEEE, 582–586. <https://doi.org/10.1109/ISIT.2019.8849834>
- [12] Adel Javanmard and Andrea Montanari. 2018. Online Rules for Control of False Discovery Rate and False Discovery Exceedance. *The Annals of Statistics* 46, 2 (2018), 526–554.
- [13] Theis Ingerslev Jensen, Bryan Kelly, and Lasse Heje Pedersen. 2023. Is There a Replication Crisis in Finance? *The Journal of Finance* 78, 5 (2023), 2465–2518.
- [14] Terence Lim, Kumar Muthuraman, and Michael Sury. 2026. QRAFT: An Agentic Framework for Empirical Research in Quantitative Finance. *arXiv preprint arXiv:2604.18500* (2026).
- [15] Aaditya Ramdas, Fanny Yang, Martin J. Wainwright, and Michael I. Jordan. 2017. Online Control of the False Discovery Rate with Decaying Memory. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [16] Daniel Russo and James Zou. 2016. Controlling Bias in Adaptive Data Analysis Using Information Theory. In *Proceedings of the 19th International Conference on Artificial Intelligence and Statistics (AISTATS)*. 1232–1240.
- [17] Meruzhan Sargsyan. 2025. Structural Enforcement of Statistical Rigor in AI-driven Discovery. *arXiv preprint* (2025).
- [18] Ryan Sullivan, Allan Timmermann, and Halbert White. 1999. Data-Snooping, Technical Trading Rule Performance, and the Bootstrap. *The Journal of Finance* 54, 5 (1999), 1647–1691. <https://doi.org/10.1111/0022-1082.00163>
- [19] Vladimir Vovk and Ruodu Wang. 2021. E-values: Calibration, Combination, and Applications. *The Annals of Statistics* 49, 3 (2021), 1736–1754.
- [20] Ruodu Wang and Aaditya Ramdas. 2022. False Discovery Rate Control with E-values. *Journal of the Royal Statistical Society: Series B* 84, 3 (2022), 822–852.
- [21] Ziyu Xu and Aaditya Ramdas. 2024. Online Multiple Testing with E-values. *arXiv preprint arXiv:2311.06412* (2024).
- [22] Yishuo Yuan, Jiayi Sheng, Sirui Zeng, Jiaqi Wang, and Jiaheng Liu. 2026. AlphaCrafter: A Full-Stack Multi-Agent Framework for Cross-Sectional Quantitative Trading. *arXiv preprint arXiv:2605.05580* (2026).